

Skeleton Code for SegNet Implementation Assignment, Computer Vision Class, ETH Zurich

Overview

- `image_segmentation/`: includes training and validation scripts.
- `lib/`: contains core functions, data preparation, model definition, and utility functions.

Installation

0. If you already followed instructions of assignment 3, then you can skip to step 4.
1. Install either [miniconda](#) or [anaconda](#).
2. Create a virtual environment and activate it:

```
conda create -n CV22_Image_Segmentation python=3.8
conda activate CV22_Image_Segmentation
```

3. Install PyTorch 1.9.1 from the [official website](#). CPU-only version is sufficient for this assignment.
4. Install dependencies.

```
pip install -r requirements.txt
```

NOTE: TensorBoard and tensorboardX may not be compatible on some platforms. If you encounter difficulties when installing them, just remove them from `requirements.txt`. They are used solely for visualizing your results, which is optional (but helpful) for building your model.

5. Add current project directory (which we will later denote as `${ROOT}`) to `PYTHONPATH` environment variable. On Linux you can do the following:

```
export PYTHONPATH=${PYTHONPATH}:${PWD}
```

Data Preparation for Multi-digit MNIST Dataset

1. Unzip the multi-digit-MNIST dataset (multi-digit-mnist-dataset.zip) to `${ROOT}`. You should have the following directory structure after unzipping:

```
${ROOT}
|-- data
    |-- multi-digit-mnist
        |-- batch00001.mat
        |...
        |-- testset001.mat
        |...
```

Implement Simplified SegNet Architecture

Complete the skeleton code in `lib/models/seg_net.py`. Please refer to the `hints in comments`, the original paper, as well as the slides from lab assignment session for details.

Training on Multi-digit MNIST Dataset

To train your model, run:

```
python image_segmentation/train_mnist.py
```

Validate the model

To validate the model after training, run:

```
python image_segmentation/validate_mnist.py
```

(Optional) If you installed TensorBoard and tensorboardX, then you should have TensorBoard logs saved to `out/logs`. You can monitor the logs (loss curves, validation visualization, etc.) on <http://localhost:6006> via:

```
tensorboard --logdir out/logs --port 6006
```

References

1. The overall structure of the code (roughly) follows [Simple Baselines for Human Pose Estimation and Tracking](#).
2. Multi-digit MNIST dataset was created using the script from [Recurrent Pixel Embedding for Instance Grouping](#).

License

This code should only be used for Computer Vision class at ETH Zurich. Please do not distribute this code outside the course.